



IMPACT DECK • HACKATHON, 1-3 OCTOBER 2026

Qaid قيد

Qaid prepares the books. The accountant decides.

An assistant for Egyptian accounting offices that turns their clients' phone-photographed receipts into reviewed, ready-to-confirm entries.

THE PROBLEM

Egypt runs on small businesses. Their books run on paper.

90%

of Egypt's private sector is micro, small and medium businesses

43%

of GDP, about EGP 1.2 trillion of production

75%

of the workforce

MSMEDA CEO, OECD committee, Nov 2024 (Ahram Online).

SOURCED

What the accountant does with it

- Every client sends receipts as photos and chat messages.
- The accountant chases the missing ones, then types each one in.
- Manual receipt entry: **3-5 minutes** each, **1-3%** errors.

56% of 250 accountants and bookkeepers surveyed say they spend too much time on manual tasks (Dext, UK, 2023).

Workload figures are UK/US and vendor sources. No Egyptian equivalent was found. **INDICATIVE**

Small businesses keep their accountant. So we serve the accountant.

Why they will not switch to software alone

- The state is adding compliance, not removing it. The simplified tax regime (Law 6 of 2025, turnover up to EGP 20m) is conditional on e-invoice and e-receipt compliance.
- The e-invoice threshold dropped from EGP 500,000 to 250,000 a year, so more small businesses must register.
- E-invoices and e-receipts cover what a business *issues*. The paper from its *suppliers* is still a photo for the accountant to read.

EY; Andersen; ClearTax Egypt. **SOURCED** The last point is our reasoning.

REASONING

Shop owner (client)

taps a link, photographs a receipt. Nothing else changes.



Qaid

reads, checks, flags, prepares




Accounting office

reviews, confirms, files. Same people, more clients.

From a phone photo to a confirmed entry

Upload returns at once; reading continues in the background.



Then: month-end Excel / PDF pack, chat over the books, Arabic reminders.

Two readers, one referee, a human on top

Who does what

AI reads text, proposes a category, drafts a reminder.

CODE adds VAT, totals, duplicates, price jumps, client status. Every figure comes from stored records.

HUMAN confirms. Only confirmed entries count.

Nothing is fixed silently. If the readers disagree, the entry is flagged.

Two lanes


A: Cohere Parse v5 reads layout, GPT-6 structures it.

B: GPT-6 vision reads the image.

fuse() compares field by field: *agree*, *resolved*, *disputed*, or *single* if one lane is down.

The Parse deployment hit its rate limit (HTTP 429) during testing. The pipeline fell back to one lane and said so.

Measured: 50 fields



agree 40
resolved 2
disputed 3
single lane 5

10 receipts, both lanes on. **MEASURED**


The client never waits. The lanes run side by side.

One receipt, seconds (measured)

One after the other



Side by side



about 9.9 s vs about 5.9 s (single-run lane times)

Measured end to end (both lanes + fuse): median 6.5 s, max 8.0 s

0s 2s 4s 6s 8s 10s

Lane times from single runs (vision 2.9–4.3 s, Parse 5.9 s). One-after-the-other is the sum, our arithmetic. Fused median 6.5 s, max 8.0 s over 10 receipts.

- **Upload returns at once.** Reading continues after the response (`after()`).
- **Postgres job queue** with `skip_locked`, retries and a dead state. A one-minute sweep is the backstop.
- **Parallel lanes.** Time is the slower lane, not the sum.
- **Live board.** It refreshes itself every 4 s, so the accountant sees the row appear.
- **Phone checks first.** Blur, light and crop are caught before upload, so bad photos never reach the models.

Not just OCR, not just one model

	PLAIN OCR	ONE LLM	QAID
Reads Arabic and English receipts	partly	yes	yes
A second, independent reading	no	no	yes, two lanes
Flags what it is unsure of	no	rarely	yes, per field
Totals and VAT computed by code	sometimes	model guess	always
Checks across time (duplicates, price jumps)	no	no	yes
Accountant has the final say	varies	varies	built in

A design comparison, not a benchmark. We did not run OCR-only or single-model baselines. QUALITATIVE

MEASURED RESULT

Money fields were perfect. We are open about the misses.

Correct fields, 10 receipts, both lanes

Subtotal		10/10
VAT		10/10
Total		10/10
Date		9/10
Vendor		9/10

Median 6.5 s per receipt, max 8.0 s.

MEASURED

WHAT WE ARE NOT CLAIMING

- Receipts were **clean renders, not phone photos**, and only **10** of them. Real accuracy will be lower.
- One date came out as 2021 instead of 2026 and one Arabic name lost a letter. Both were chosen between two disagreeing readers, so the next fix is to send every such pick to the accountant.
- The Parse reader was rate limited, so 5 of 50 fields were single-lane. Production needs more quota.

More clients per bookkeeper, same team

If Qaid takes a client from **5 h** to **4 h** a month, a bookkeeper moves from 25 to about **31 clients**: +6.

Egyptian bookkeeper pay is about EGP 6,000 a month (Glassdoor median), so an hour costs about EGP 37.5. The 5 h and 4 h figures are US workload numbers.

Pricing idea: **EGP 25-35 per client per month**, or a flat office plan, until a real fee is known.

ALL ESTIMATES

Extra monthly gain per bookkeeper, after paying EGP 25 per client

OFFICE FEE PER CLIENT	GAIN (EGP)
375 (2.0x labour)	+1,563
469 (2.5x)	+2,148
563 (3.0x)	+2,734

Fees come from an assumed 2-3x markup on labour, not from Egyptian office data. We have not found a real fee per client. ESTIMATE

What Qaid does not do, and what we do next

- Prepares only. It does not file taxes, talk to banks or send WhatsApp messages.
- Printed receipts, EGP only. Handwriting is out of scope.
- VAT and insurance rates in config are still marked "to verify".
- Test with one real accounting office and real phone photos.
- Flag every disputed or resolved field for review.
- Raise the Parse quota, then pilot with a handful of clients.

Qaid prepares. The accountant decides.